

# Set Up

In converting the implementations to use ExecutorService, I created 4 different implementations:

- Single Thread (TCPEchoServerSingleExecutor)
- Thread Per Connection (TCPEchoServerThread)
- Cached Thread Pool (TCPEchoServerCachedPoolExecutor)
- Fixed Thread Pool (TCPEchoServerFixedPoolExecutor)

Upon comiling (mvn clean compile), I ran each and assigned them to their own port, so that I could run my Jmeter tests en masse. The following commands will run each of the implementations:

```
-java -cp target/classes TCPEchoServerSingleExecutor 8081
```

```
-java -cp target/classes TCPEchoServerThread 8082
```

```
-java -cp target/classes TCPEchoServerCachedPoolExecutor 8083
```

```
-java -cp target/classes TCPEchoServerFixedPoolExecutor 8084 50 (or 10, this argument is the thread pool size)
```

I also created an EchoProtocol class to handle receipt and echo of the messages, with delay when message size  $\geq 4$

```
INFO: Echoed 16 bytes to /127.0.0.1:51444
Jul 07, 2025 7:11:26 PM EchoProtocol handleEchoClient
INFO: Handling client at /127.0.0.1:51445
Jul 07, 2025 7:11:26 PM EchoProtocol handleEchoClient
INFO: Received 16 bytes from /127.0.0.1:51445
Jul 07, 2025 7:11:26 PM EchoProtocol handleEchoClient
INFO: Echoed 16 bytes to /127.0.0.1:51445
PS C:\Users\nickb\IdeaProjects\Flyweight> java -cp target/classes TCPEchoServerSingleExecutor 8081
```

```
INFO: Echoed 16 bytes to /127.0.0.1:52116
Jul 07, 2025 7:12:12 PM EchoProtocol handleEchoClient
INFO: Echoed 16 bytes to /127.0.0.1:52091
Jul 07, 2025 7:12:12 PM EchoProtocol handleEchoClient
INFO: Echoed 16 bytes to /127.0.0.1:52093
PS C:\Users\nickb\IdeaProjects\Flyweight> java -cp target/classes TCPEchoServerThread 8082
```

```
INFO: Echoed 16 bytes to /127.0.0.1:53132
Jul 07, 2025 7:17:08 PM EchoProtocol handleEchoClient
INFO: Echoed 16 bytes to /127.0.0.1:53130
Jul 07, 2025 7:17:08 PM EchoProtocol handleEchoClient
INFO: Echoed 16 bytes to /127.0.0.1:53131
PS C:\Users\nickb\IdeaProjects\Flyweight> java -cp target/classes TCPEchoServerCachedPoolExecutor 8083
```

```
Jul 07, 2025 7:38:20 PM EchoProtocol handleEchoClient
INFO: Echoed 16 bytes to /127.0.0.1:63234
Jul 07, 2025 8:03:04 PM TCPEchoServerFixedPoolExecutor lambda$main$0
INFO: Shutting down executor...
PS C:\Users\nickb\IdeaProjects\Flyweight> java -cp target/classes TCPEchoServerFixedPoolExecutor 8084 50
```

# Cached vs Fixed Thread Pool vs Thread Per Connection

The cached thread pool from `ExecutorService` allows for creation of new threads as new client requests are made (so lazy initialization). However, where it stands out from the other implementations is that the threads are cached and thus can be re-used when available. This is what makes it a great example of the Flyweight design pattern, since Flyweight uses a factory to create new objects when needed but reuse existing if it already exists. Fixed Thread Pool, on the other hand sets a predetermined pool size of threads and appropriates those as requests come in. Once all threads are used, no further all requests must wait until a thread is available. As I show a little later, Cached Thread Pool yields better performance.

Thread Per Connection has similarities to cached thread pool in that new threads are created as new connections come in. However, thread per connection does not reuse threads from a cache like cached thread pool does. Since there is not a direct similarity in `ExecutorService` for thread per connection, I used the original class to show how it compares performance-wise to the other implementations.


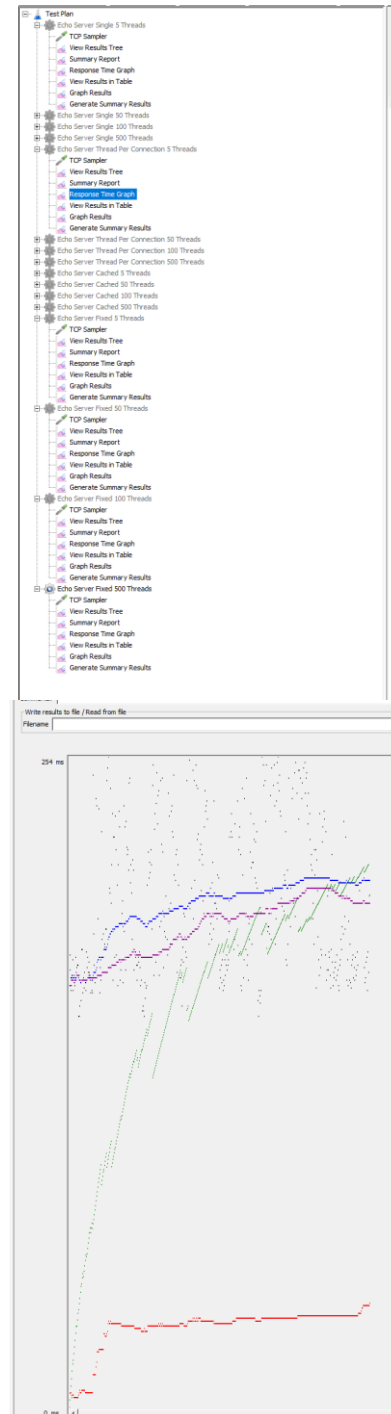
I also have a single thread implementation. As you will see, this is the worst of the group in terms of performance by a large margin.

# Jmeter & Power BI

After starting each of my own implantations on their own port, I used Jmeter to simulate 5, 50, 100, and 500 requests to each of them. I ran the simulation for fixed thread pool twice, once with a pool size of 10 and once with a pool size of 50.

Jmeter has its own suite of visuals (as you can see in the examples to the right). However, I did not think they were good at telling the visual story about which implementation performed the best. Thus, I exported the results of the tests to create my own charts in Power BI to evaluate performance.

The following slides show my results for each connection type with 5, 50, 100, and 500 threads

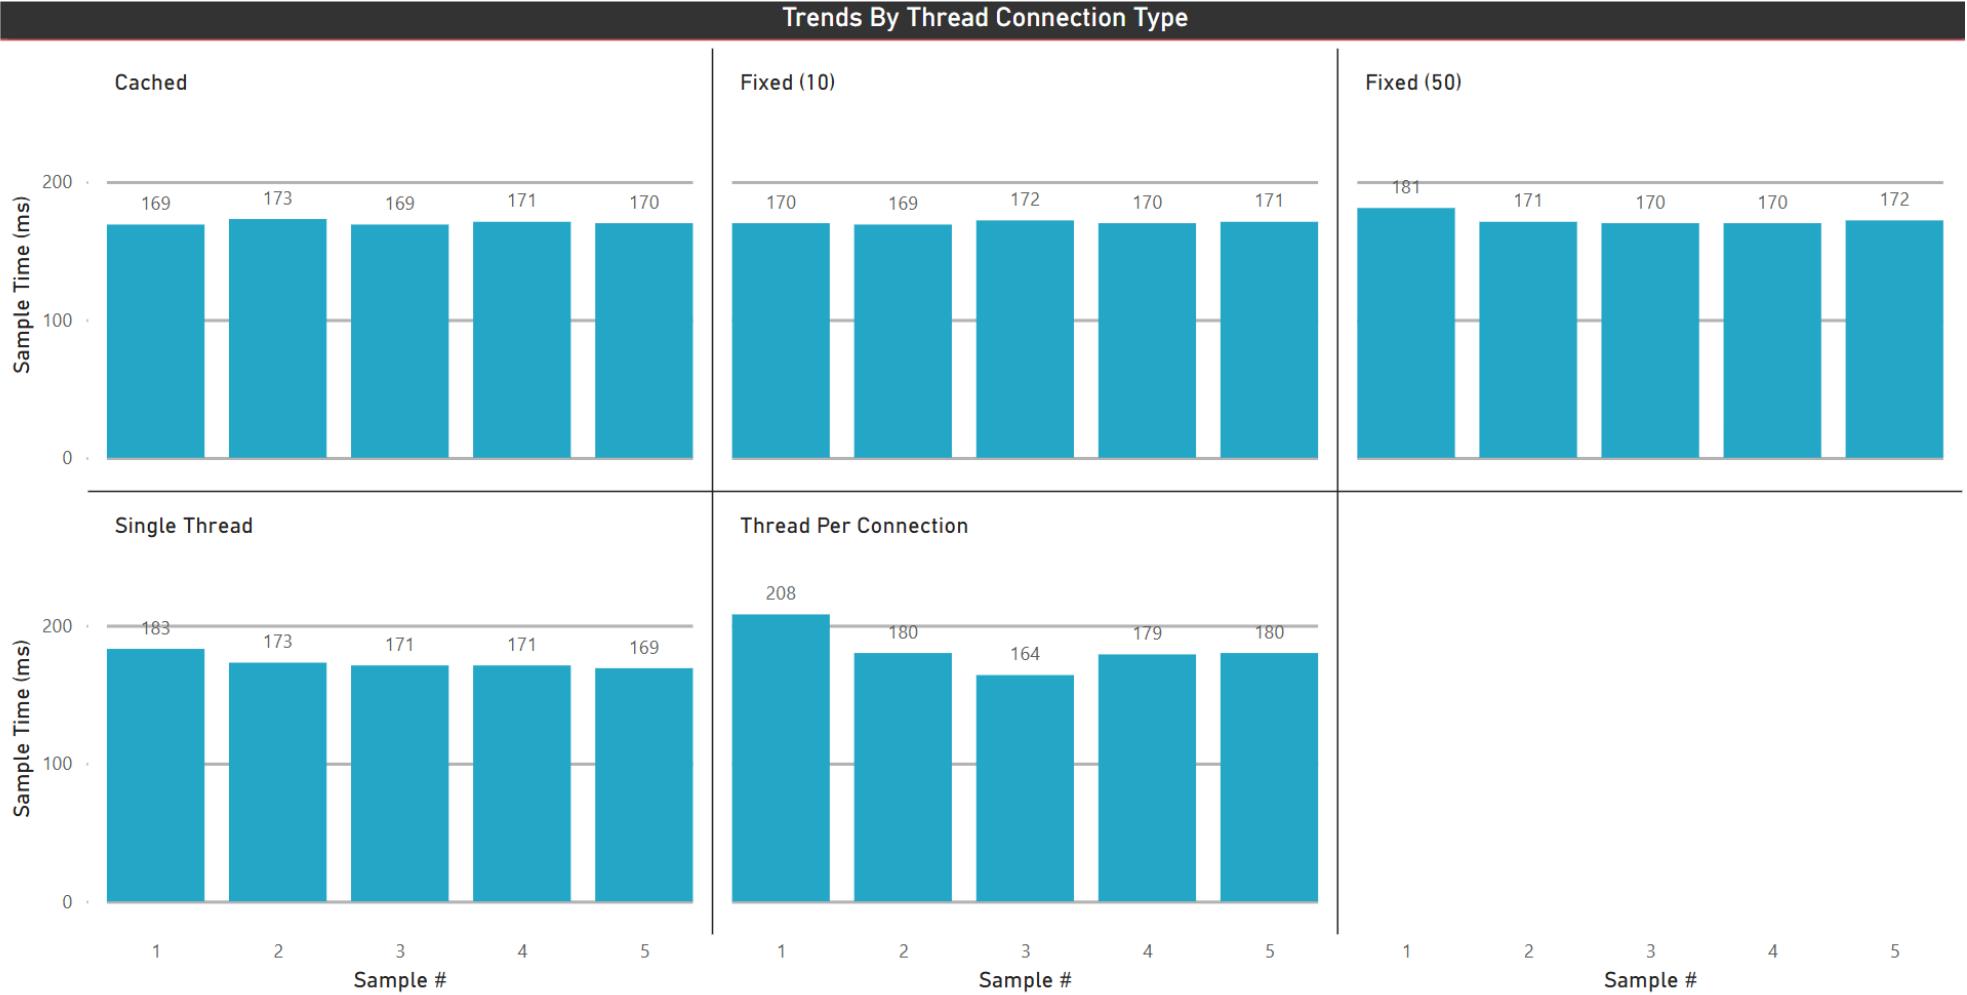


Text	Sampler result	Request	Response data
✓ TCP Sampler	Thread Name:Echo Server Thread Per Connection 500 Threads 1-9		
✓ TCP Sampler	Sample Start:2025-07-07 19:12:11 EDT		
✓ TCP Sampler	Load time:176		
✓ TCP Sampler	Connect Time:1		
✓ TCP Sampler	Latency:115		
✓ TCP Sampler	Size in bytes:16		
✓ TCP Sampler	Sent bytes:0		
✓ TCP Sampler	Headers size in bytes:0		
✓ TCP Sampler	Body size in bytes:16		
✓ TCP Sampler	Sample Count:1		
✓ TCP Sampler	Error Count:0		
✓ TCP Sampler	Data type ("text" "bin" ""):text		
✓ TCP Sampler	Response code:200		
✓ TCP Sampler	Response message:OK		
✓ TCP Sampler	SampleResult fields:		
✓ TCP Sampler	ContentType:		
✓ TCP Sampler	DataEncoding: windows-1252		

# 5 Threads

With only 5 threads, no implementation stood out with each performing relatively similarly

## 5 Threads




# 50 Threads

When the load is increased to 50 threads, the Sample Time of the Single Thread implementation skyrockets.

This makes sense as there is only one thread that must be waited on by each request.

The other 4 implementations stay relatively consistent throughout the run.

## 50 Threads



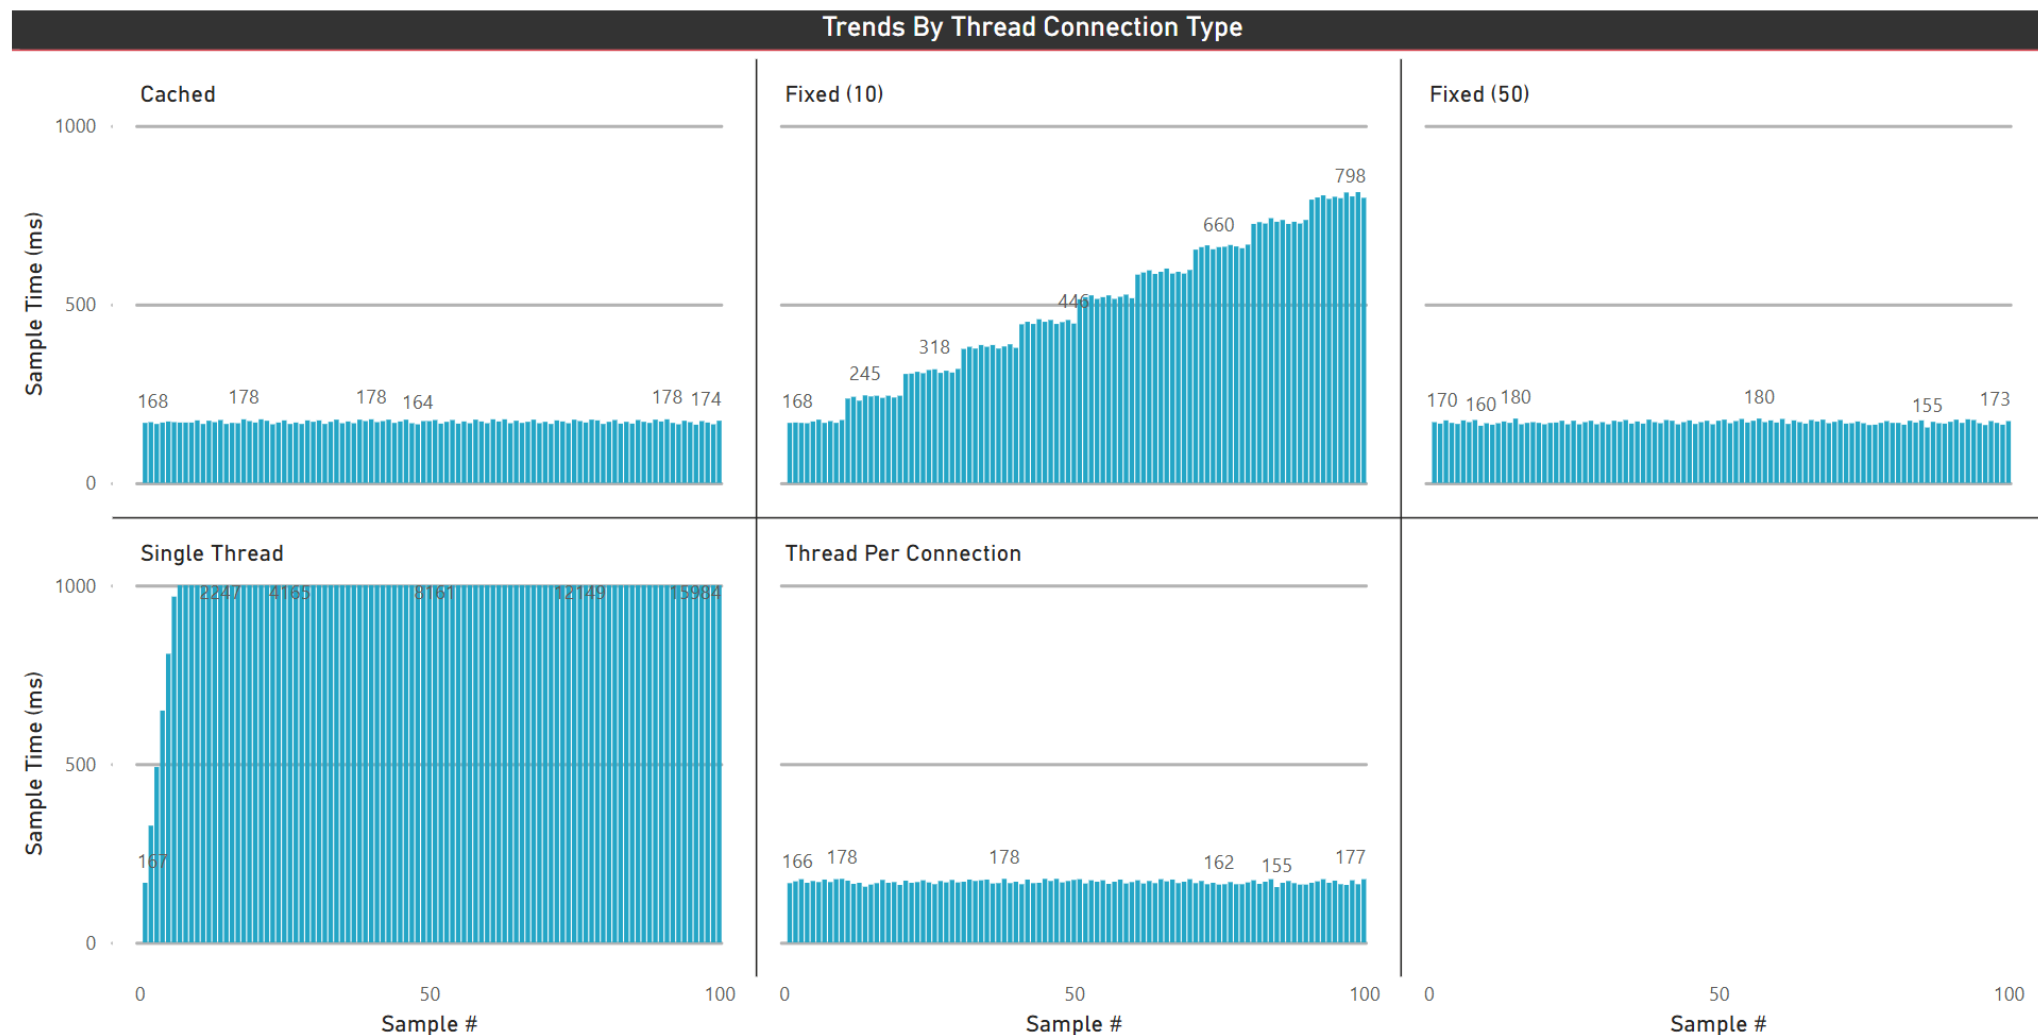
# 100 Threads

As the load increases, weak points in certain implementations begin to show themselves. Single Thread continues to balloon out of control while the Fixed Thread Pool with 10 pool size begins to suffer.

This makes sense, since there is only a pool of 10 threads to go around, a load of 100 threads is going to have many requests waiting.

Cached, Fixed (50 pool size), and Thread Per Connection continue to remain consistent.

## 100 Threads




# 500 Threads

The 500 thread load showed the cached thread pool is the best implementation for performance. Single Thread hit the response time cap and began to fail. The Fixed Thread Pool (10) continued to degrade while the 50 pool size began its degradation (as expected).

The Thread Per Connection also began to suffer, albeit no where near as bad as the previous 3. It can be seen that there are much higher times at the end of the run than at the beginning whereas the cached implementation stays consistent throughout. This is likely due to the overhead Thread Per Connection has from creating a new thread for every connection where as cached can reuse threads from its cache.

The following slide shows an average (rather than time series) of the 500 thread response times to show this point even better.

## 500 Threads



# 500 Threads (Average Time)

## 500 Threads

